

Computer Vision

CVI620

Session 22
03/2025

What is Left?

4 sessions

1. Optimization and Loss Function
2. Code + Logistic Regression
3. ML and Images
4. Perceptron and Neural Networks
5. Deep Neural Networks
6. Convolution Neural Networks (CNN)
7. Advanced CNNs
8. Project
9. Segmentation
10. Introduction to object detection and image generation methods with AI
11. Project

Agenda

Review


01

Convolutional
Neural Networks

02

Code

02





Review

Data preprocessing


- Resize
- Normalize
- Flatten
- Create labels
- Label encoding
- One-hot encoding
- Train-test split



Review

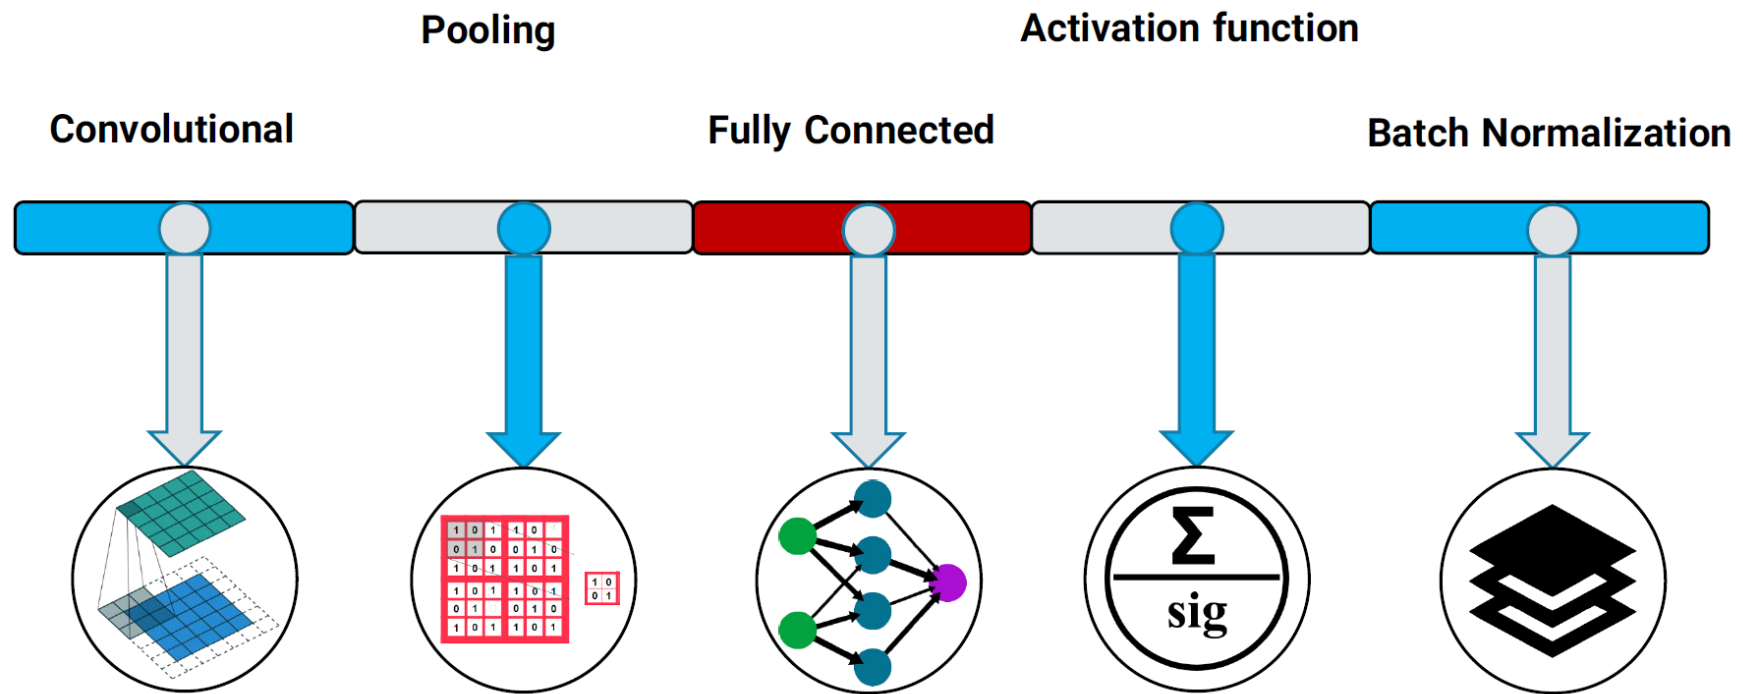
Model

- Sequential
- Layers.Dense
- Activation: step, sigmoid, softmax, relu
- Optimzier: SGD, adam
- Loss: MAE, MSE, categorical_crossentropy
- batch_size
- epoch




Convolutional Neural Networks (CNNs)


Layers in CNNs




Convolution Layer



Multiple Convolutions



Convolution for Colored Images



Padding

0	0	0	0	0	0	0	0
0	3	3	4	4	7	0	0
0	9	7	6	5	8	2	0
0	6	5	5	6	9	2	0
0	7	1	3	2	7	8	0
0	0	3	7	1	8	3	0
0	4	0	4	3	2	2	0
0	0	0	0	0	0	0	0

$6 \times 6 \rightarrow 8 \times 8$

*

1	0	-1
1	0	-1
1	0	-1

3×3


=

-10	-13	1			
-9	3	0			


6×6

Stride


S = 1



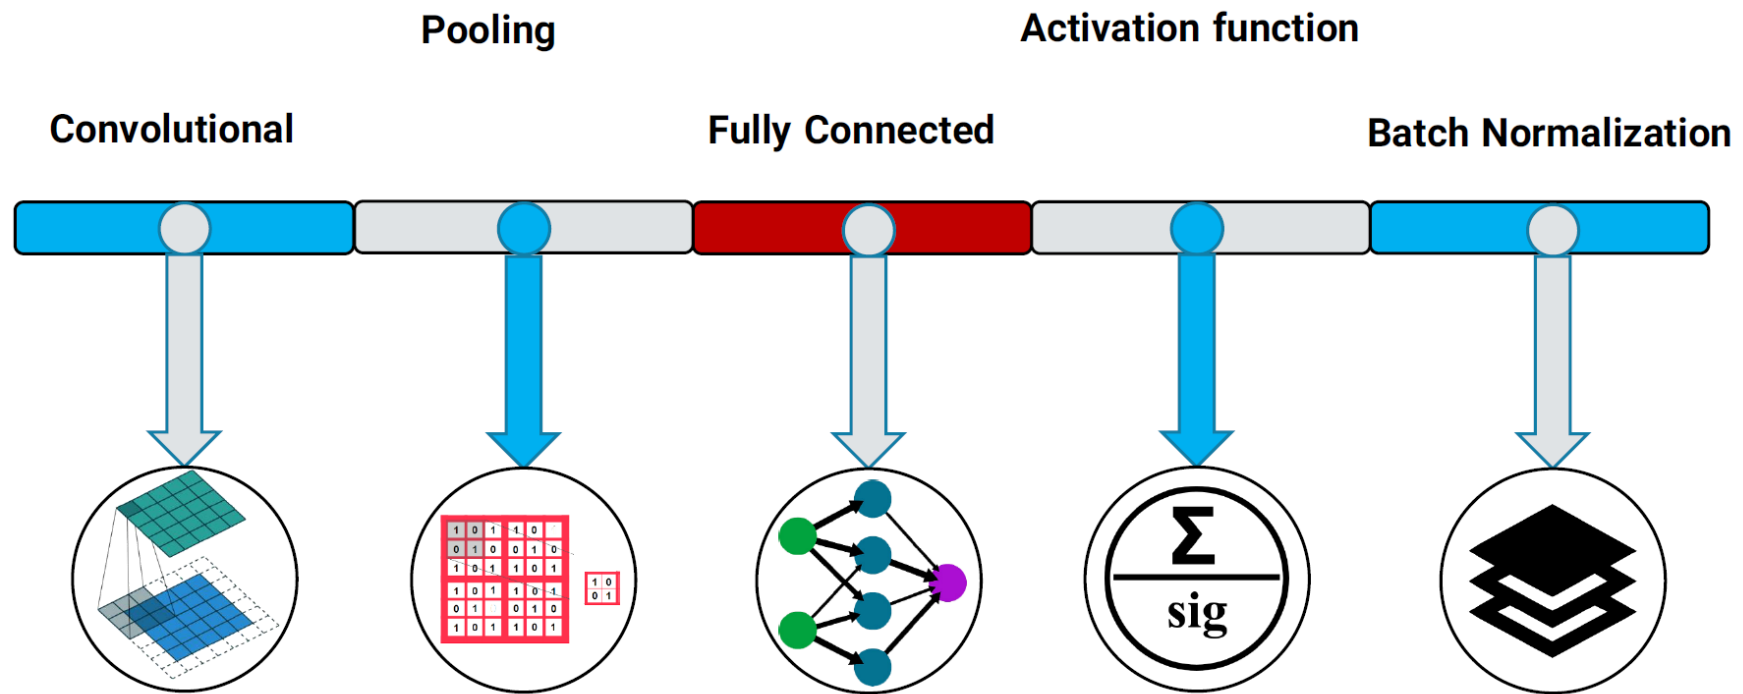
S = 3



Post Convolution Dimensions



Layers in CNNs



MaxPooling

2	2	7	3
9	4	6	1
8	5	2	4
3	1	2	6

Max Pool
→
Filter - (2 x 2)
Stride - (2, 2)

9	7
8	6

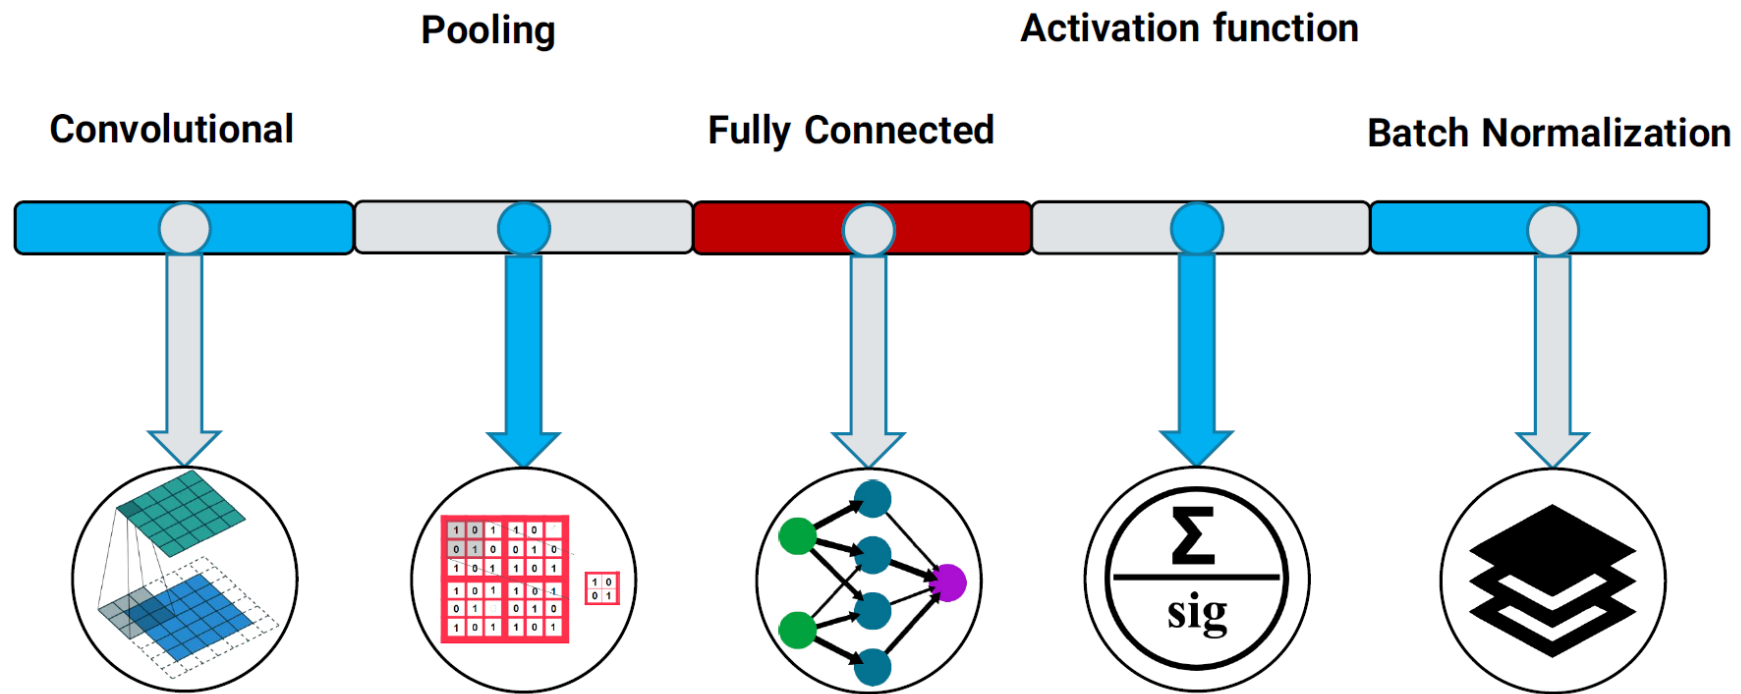
Average Pooling

2	2	7	3
9	4	6	1
8	5	2	4
3	1	2	6


Average Pool
→
Filter - (2 x 2)
Stride - (2, 2)

4.25	4.25
4.25	3.5


Layers in CNNs



MLP



CNNs



Conv2D

```
tf.keras.layers.Conv2D(  
    filters, kernel_size, strides=(1, 1), padding='valid', data_format=None,  
    dilation_rate=(1, 1), groups=1, activation=None, use_bias=True,  
    kernel_initializer='glorot_uniform', bias_initializer='zeros',  
    kernel_regularizer=None, bias_regularizer=None, activity_regularizer=None,  
    kernel_constraint=None, bias_constraint=None, **kwargs  
)
```


```
net = models.Sequential([
    layers.Conv2D(32, (3,3), activation='relu', input_shape=(32,32,3)),
    layers.MaxPool2D(),
    layers.Conv2D(32, (3,3), activation='relu'),
    layers.MaxPool2D(),
    layers.Flatten(),
    layers.Dense(100, activation='relu'),
    layers.Dense(2, activation='softmax')
])
```

```
net.compile(optimizer='SGD',
            loss='categorical_crossentropy',
            metrics='accuracy')
```

```
H = net.fit(X_train, y_train, validation_data=(X_test, y_test), batch_size=64, epochs=10)
```

```
net.save('S22/CNN_classifier.h5')
```

Overfitting



- Model works good on train data but not on test data.
- It memorizes and not generalize

Batch Normalization

We normalized the input data.

Why not normalizing values after each layer?

A technique to normalize activations in a neural network, layer by layer, during training.

- Stabilizes learning
- Speeds up convergence
- Reduces internal covariate shift
- Allows higher learning rates



Data Augmentation

- A technique to increase the size and diversity of training data by applying random transformations.
- More robust models with better performance on unseen data.

Rotation



Width Shift



Brightness



Shear



Zoom



ImageDataGenerator

- A tool to apply data augmentation

```
aug = ImageDataGenerator(rotation_range=30,  
                          width_shift_range=0.1,  
                          height_shift_range=0.1,  
                          shear_range=0.2,  
                          zoom_range=0.2,  
                          horizontal_flip=True,  
                          fill_mode="nearest")
```

Learning Rate


A hyperparameter that controls how much the model updates its weights in response to the loss gradient.

- Too high → Model diverges
- Too low → Slow convergence
- Just right → Fast & stable training


```
from keras.optimizers import SGD
...
opt=SGD(lr=..., decay=...)
net.compile(..., optimizer=opt)
```


100

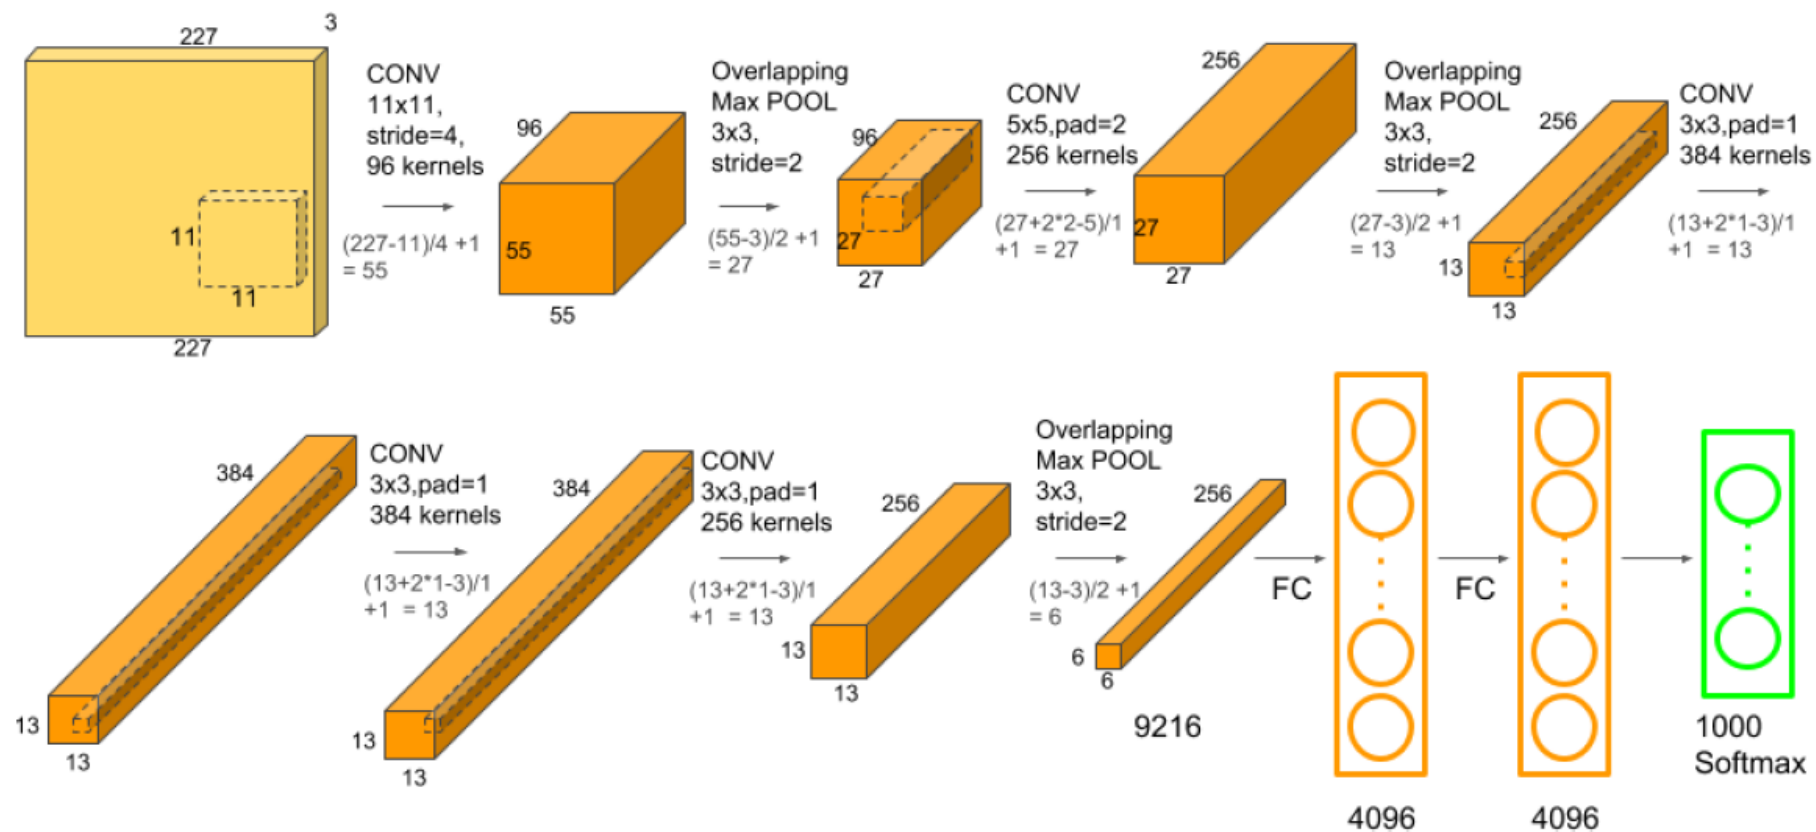
- ImageNet dataset
- 1.2 m images with 1000 classes
- Classification
- Single object detection
- Object detection



Results



AlexNet



VGG

